

Class4_NN

January 29, 2025

```
[1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split

import torch
from torch import nn
from torch import optim
from torch.utils.data import Dataset, DataLoader
```

```
[2]: device = 'cuda' if torch.cuda.is_available() else 'cpu'
print(device)
```

cuda

0.0.1 Dataset preparation

- Import Dataset and Get Properties

```
[3]: df = pd.read_csv(r'./dataset/WineQT.csv')

target_name = 'quality' # integer between 0 - 10

print("Datapoint shape:", df.shape)
print("Attribute:", list(df.columns))
```

Datapoint shape: (1143, 13)

Attribute: ['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar', 'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH', 'sulphates', 'alcohol', 'quality', 'Id']

```
[4]: # Remove unnecessary attribute
df = df.drop(columns=["Id"])
```

- Inspect dataset

```
[5]: stat = df.describe()
mean, std = stat.iloc[1], stat.iloc[2]
display(pd.concat([mean, std], axis=1).T)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
mean	8.311111	0.531339	0.268364	2.532152	0.086933	
std	1.747595	0.179633	0.196686	1.355917	0.047267	

	free sulfur dioxide	total sulfur dioxide	density	pH	\
mean	15.615486	45.914698	0.996730	3.311015	
std	10.250486	32.782130	0.001925	0.156664	

	sulphates	alcohol	quality
mean	0.657708	10.442111	5.657043
std	0.170399	1.082196	0.805824

- Preprocessing

Z transform (Normalization)

```
[6]: def z_norm(dataset: pd.DataFrame, target: str):
    answer = dataset[target]
    dataset = dataset.drop(columns=[target])
    for column in dataset.columns:
        dataset[column] = (dataset[column] - dataset[column].mean()) /
        dataset[column].std()
    display(dataset.describe()[1:3].astype(np.int32))

    return dataset, answer.to_numpy(dtype=np.int8)

train, test = z_norm(df.copy(), target_name)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
mean	0	0	0	0	0	
std	1	1	1	0	0	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
mean	0	0	0	0	0	
std	0	1	0	0	0	

	alcohol
mean	0
std	1

Train test split

```
[7]: import random
    random.seed(42)
    np.random.seed(42)
    x_train, x_test, y_train, y_test = train_test_split(train, test, train_size=0.
    8, shuffle=True)
    print(len(x_train), len(y_test))
```

914 229

Convert to torch.Tensor

```
[8]: x_train, x_test, y_train, y_test = (
    torch.tensor(x_train.to_numpy(), dtype=torch.float32),
    torch.tensor(x_test.to_numpy(), dtype=torch.float32),
    torch.tensor(y_train, dtype=torch.long),
    torch.tensor(y_test, dtype=torch.long),
)

# Remove offset
y_train = y_train - y_train.min()
y_test = y_test - y_test.min()
```

```
[9]: print("train:", x_train.shape)
```

train: torch.Size([914, 11])

0.0.2 Pipeline setup

```
[10]: # Must Preprocess first
class WineDataset(Dataset):
    def __init__(self, x, y):
        self.features = x
        self.target = y

    def __len__(self):
        return len(self.features)

    def __getitem__(self, idx):
        return self.features[idx].to(device), self.target[idx].to(device)

train_dataset = WineDataset(x_train, y_train)
test_dataset = WineDataset(x_test, y_test)

train_loader = DataLoader(train_dataset, batch_size=32)
test_loader = DataLoader(test_dataset, batch_size=32)
```

0.0.3 Model Setup

- Define Model

```
[11]: class RegressionNN(nn.Module):
    def __init__(self, in_dim, out_dim, dtype):
        super().__init__()
        self.input_layer = nn.Linear(in_dim, 64, dtype=dtype)
        self.hidden_layer = nn.Linear(64, 32, dtype=dtype)
        self.output_layer = nn.Linear(32, out_dim, dtype=dtype)
```

```

        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.input_layer(x))
        x = self.relu(self.hidden_layer(x))
        y = self.output_layer(x)
        return y

```

```
[12]: next(iter(train_loader))[0].shape, next(iter(train_loader))[1].shape
```

```
[12]: (torch.Size([32, 11]), torch.Size([32]))
```

```
[13]: len(np.unique(y_train.numpy()))
```

```
[13]: 6
```

```
[14]: torch.manual_seed(42)

model = RegressionNN(11, len(np.unique(y_train.numpy())), torch.float32).
    ↪to(device)

loss_fn = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), 1e-3)
```

0.0.4 Training Loop

```
[15]: model.train()
for i in range(1, 101):
    iter_loss = 0
    for batch in train_loader:
        input_batch, target_batch = batch

        optimizer.zero_grad()

        outputs = model(input_batch)

        loss = loss_fn(outputs, target_batch)
        loss.backward()

        optimizer.step()

        iter_loss += loss

    if i%20 == 0:
        print(f"loss at iter {i} = {loss/len(train_loader)}")
```

```
loss at iter 20 = 0.03160659968852997
loss at iter 40 = 0.026845039799809456
loss at iter 60 = 0.02331145852804184
loss at iter 80 = 0.020234109833836555
loss at iter 100 = 0.018250463530421257
```

0.0.5 Eval Model

```
[16]: model.eval()

total_loss = 0

for i, batch in enumerate(test_loader):
    input_batch, target_batch = batch

    with torch.no_grad():
        outputs = model(input_batch)
        loss = loss_fn(outputs, target_batch)
        total_loss += loss

print(f"loss of testset = {total_loss/len(test_loader)}")
```

```
loss of testset = 3.3306384086608887
```