

Class5_NN

February 4, 2025

```
[1]: import numpy as np

import torch
from torch import nn
from torch import optim
from torch.utils.data import Dataset, DataLoader
from torchvision import datasets, transforms
```

```
[2]: device = 'cuda' if torch.cuda.is_available() else 'cpu'
print(device)
```

cuda

0.0.1 Dataset preparation

- Import Dataset and Get Properties

```
[12]: training_data = datasets.FashionMNIST(root="./dataset", train=True,
↳download=False, transform=transforms.ToTensor())
testing_data = datasets.FashionMNIST(root="./dataset", train=False,
↳download=False, transform=transforms.ToTensor())
print(training_data, "\n#####\n", testing_data)
```

Dataset FashionMNIST

Number of datapoints: 60000

Root location: ./dataset

Split: Train

StandardTransform

Transform: ToTensor()

#####

Dataset FashionMNIST

Number of datapoints: 10000

Root location: ./dataset

Split: Test

StandardTransform

Transform: ToTensor()

- Inspect dataset

```
[14]: print("training:", training_data.data.size())
      print("testing:", testing_data.data.size())
```

```
training: torch.Size([60000, 28, 28])
testing: torch.Size([10000, 28, 28])
```

```
[17]: training_data.targets
```

```
[17]: tensor([9, 0, 0, ..., 3, 0, 5])
```

0.0.2 Pipeline setup

```
[69]: # Must Preprocess first
class ClassificationDataset(Dataset):
    def __init__(self, x, y):
        self.features = x.to(torch.float)
        self.target = y.to(torch.float)

    def __len__(self):
        return len(self.features)

    def __getitem__(self, idx):
        return self.features[idx].to(device), self.target[idx].to(device)

train_dataset = ClassificationDataset(training_data.data, training_data.targets)
test_dataset = ClassificationDataset(testing_data.data, testing_data.targets)

train_loader = DataLoader(train_dataset, batch_size=256)
test_loader = DataLoader(test_dataset, batch_size=256)
```

0.0.3 Model Setup

- Define Model

```
[64]: class ClassificationNN(nn.Module):
    def __init__(self, out_dim, dtype=torch.float32):
        super().__init__()
        self.extract_layer = nn.Conv2d(1, 16, 1)
        self.bn1 = nn.BatchNorm2d(16)
        self.hidden_layer = nn.Conv2d(16, 16, 3)
        self.bn2 = nn.BatchNorm2d(16)
        self.expose_layer = nn.Conv2d(16, 1, 1)
        self.bn3 = nn.BatchNorm2d(1)

        self.linear_layer = nn.Linear(1* 26* 26, out_dim)

    def forward(self, x):
```

```

        x = self.bn1(self.extract_layer(x))
        x = self.bn2(self.hidden_layer(x))
        x = self.bn3(self.expose_layer(x))

        # Logit
        y = self.linear_layer(torch.flatten(x, start_dim=1, end_dim=-1))

        return y

    def prediction(self, logit):
        with torch.no_grad():
            prob = nn.functional.softmax(logit, dim=1)
            outputs = torch.argmax(prob, dim=1)
            return outputs

```

```

[70]: torch.manual_seed(42)

model = ClassificationNN(len(training_data.classes)).to(device)

loss_fn = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), 1e-3)

```

0.0.4 Training Loop

```

[71]: model.train()
for i in range(1, 101):
    last_iter_loss = 0
    for batch in train_loader:
        input_batch, target_batch = batch

        optimizer.zero_grad()

        outputs = model(input_batch.unsqueeze(1))

        loss = loss_fn(outputs, target_batch.to(torch.long))
        loss.backward()

        optimizer.step()

        last_iter_loss = loss

    if i%20 == 0:
        print(f"loss at iter {i} = {last_iter_loss}")

```

```

loss at iter 20 = 0.30945494771003723
loss at iter 40 = 0.2895567715167999
loss at iter 60 = 0.28571146726608276

```

```
loss at iter 80 = 0.2837536931037903
loss at iter 100 = 0.28216496109962463
```

0.0.5 Eval Model

```
[76]: from sklearn.metrics import accuracy_score, confusion_matrix
import seaborn as sns
```

```
[77]: model.eval()

total_loss = 0

y_true = []
y_pred = []

for i, batch in enumerate(test_loader):
    input_batch, target_batch = batch

    with torch.no_grad():
        outputs = model(input_batch.unsqueeze(1))
        loss = loss_fn(outputs, target_batch.to(torch.long))
        total_loss += loss


    y_true.extend(target_batch.cpu().numpy())
    y_pred.extend(model.prediction(outputs).detach().cpu().numpy())

print(f"loss of testset = {total_loss/len(test_loader)}")
print(f"accuracy = {accuracy_score(y_true, y_pred)}")
```

```
loss of testset = 0.4844067692756653
accuracy = 0.8335
```

```
[79]: import matplotlib.pyplot as plt
```

```
[82]: conf_matrix = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, xticklabels=range(10), yticklabels=range(10), cmap="Blues")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.show()
```



[]: